

Project: Wrangling and Analyze Data

Data Gathering

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import requests
import os
from PIL import Image
from io import BytesIO
import json

weratedogs_archive = pd.read_csv('twitter-archive-enhanced.csv')
```

In [3]:

```
#Creating a folder if it doesn't already exist

folder_name = 'image_prediction'
if not os.path.exists(folder_name):
    os.makedirs(folder_name)
```

In [4]:

```
#Programmatic download of the TSV file

url = 'https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-predictions/image-predictions.tsv'
response = requests.get(url)
```

In [5]:

```
with open(os.path.join(folder_name, url.split('/')[-1]), mode = 'wb') as file:
    file.write(response.content)
```

In [5]:

```
image_prediction = pd.read_csv('/home/workspace/image_prediction/image-predictions.tsv', sep='\t')
image_prediction
```

1. Use the Tweepy library to query additional data via the Twitter API (tweet_json.txt)

In [7]:

```
file1 = open('tweet-json.txt', 'r')
tweet_json = file1.readlines()
```

In [8]:

```
tweets_converted = []
for tweet in tweet_json:
    tweets_converted.append(json.loads(tweet))
```

In [9]:

```
#tweet ID, retweet count, and favorite count."
df_list = []
for tweet in tweets_converted:
    tweet_id = int(tweet['id_str'])
    retweets = tweet['retweet_count']
    number_of_likes = tweet['favorite_count']
    # Append to list of dictionaries
    df_list.append({'tweet_id': tweet_id,
                    'retweets': retweets,
                    'number_of_likes': number_of_likes})
```

In [6]:

```
# Create DataFrame from list of dictionaries
tweets = pd.DataFrame(df_list, columns = ['tweet_id', 'retweets', 'number_of_likes'])
tweets
```

Assessing Data

In [11]:

```
#Programmatic Assessment
weratedogs_archive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id    78 non-null float64
in_reply_to_user_id      78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id      181 non-null float64
retweeted_status_user_id  181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls            2297 non-null object
rating_numerator          2356 non-null int64
rating_denominator        2356 non-null int64
name                     2356 non-null object
doggo                     2356 non-null object
floofer                  2356 non-null object
pupper                   2356 non-null object
puppo                     2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB
```

In [12]:

```
#Programmatic Assessment  
weratedogs_archive.duplicated().sum()
```

Out[12]:

0

In [13]:

```
#Programmatic Assessment  
image_prediction.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 2075 entries, 0 to 2074  
Data columns (total 12 columns):  
tweet_id      2075 non-null int64  
jpg_url       2075 non-null object  
img_num       2075 non-null int64  
p1            2075 non-null object  
p1_conf       2075 non-null float64  
p1_dog        2075 non-null bool  
p2            2075 non-null object  
p2_conf       2075 non-null float64  
p2_dog        2075 non-null bool  
p3            2075 non-null object  
p3_conf       2075 non-null float64  
p3_dog        2075 non-null bool  
dtypes: bool(3), float64(3), int64(2), object(4)  
memory usage: 152.1+ KB
```

In [14]:

```
#Programmatic Assessment  
image_prediction.duplicated().sum()
```

Out[14]:

0

In [15]:

```
#Programmatic Assessment  
tweets.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 2354 entries, 0 to 2353  
Data columns (total 3 columns):  
tweet_id      2354 non-null int64  
retweets      2354 non-null int64  
number_of_likes 2354 non-null int64  
dtypes: int64(3)  
memory usage: 55.2 KB
```

In [16]:

```
tweets.duplicated().sum()
```

Out[16]:

0

Quality issues

weratedogs_archive columns:

Visual Assessment

1. **name**: this column shows some unrealistic name type like single letters(a, the), there are also missing names as well as duplicated names, some name are proper case while some are all lowercase as observed via visual assessment.

-
1. **retweeted_status_id | retweeted_status_user_id | retweeted_status_timestamp**: missing entries for the following variables/columns.

-
1. **doggo | floofer | pupper | puppo**: missing entries for the following columns.

-
1. **in_reply_to_status_id | in_reply_to_user_id**: Observance of missing values from these columns.

-
1. **floofer**: Column name error. Floofer is not a dog stage according to the Dogtionary.

Programmatic Assessment

1. **expanded_urls**: during programmatic assessment, some values appear to be missing from this column.

-
1. **timestamp**: data type in this column is wrong. Should be a datetime format data type.

-
1. **tweet_id**: data type in this column is not preffered. Since I won't be using the figures here to perform any calculations, it is best practice that it is covertred to a string.
-

Tidiness issues

weratedogs_archive columns:

1. **doggo | floofer | pupper | puppo**: The following columns violate the first rule of tidiness: that each variable forms a column. They all belong under one variable: stage.

-
1. **retweeted_status_id | retweeted_status_user_id | retweeted_status_timestamp** : The following columns are not needed for the current process as majority of the data within are missing.

Cleaning Data

In [17]:

```
# Make copies of original pieces of data
weratedogs_archive_clean = weratedogs_archive.copy()
tweets_clean = tweets.copy()
image_prediction_clean = image_prediction.copy()
```

Tidiness:

Define: I will be getting rid of the unwanted columns by using the `.drop()` method.

Code

In [18]:

```
#This code drops all columns that are either unwanted from the `weratedogs_archive_clean` dataframe

weratedogs_archive_clean = weratedogs_archive_clean.drop(['retweeted_status_id', 'retweeted_status_user_id', 'retweeted_status_timestamp', 'expanded_urls', 'floofer', 'in_reply_to_status_id', 'in_reply_to_user_id' ], axis=1)
```

Test

In [19]:

```
#checking if all unwanted columns have been successfully removed from the 'weratedogs_archive_clean' dataframe
```

```
weratedogs_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 10 columns):
tweet_id          2356 non-null int64
timestamp         2356 non-null object
source            2356 non-null object
text              2356 non-null object
rating_numerator  2356 non-null int64
rating_denominator 2356 non-null int64
name              2356 non-null object
doggo             2356 non-null object
pupper           2356 non-null object
puppo            2356 non-null object
dtypes: int64(3), object(7)
memory usage: 184.1+ KB
```

Re-Gathering:

Define: I will be re-gathering the dog stages, rating_numerator , rating_denominator and name columns using the .extract() method.

Code

In [20]:

```
weratedogs_archive_clean['doggo'] = weratedogs_archive_clean.text.str.extract('(doggo)', expand = True)
weratedogs_archive_clean['pupper'] = weratedogs_archive_clean.text.str.extract('(pupper)', expand = True)
weratedogs_archive_clean['puppo'] = weratedogs_archive_clean.text.str.extract('(puppo)', expand = True)
weratedogs_archive_clean['blep'] = weratedogs_archive_clean.text.str.extract('(blep)', expand = True)
weratedogs_archive_clean['floof'] = weratedogs_archive_clean.text.str.extract('(floof)', expand = True)
weratedogs_archive_clean['snoot'] = weratedogs_archive_clean.text.str.extract('(snoot)', expand = True)
weratedogs_archive_clean['rating_numerator'] = weratedogs_archive_clean.text.str.extract('(\d+\S?\d+)(/\d+)', expand = True)
weratedogs_archive_clean['rating_denominator'] = weratedogs_archive_clean.text.str.extract('( /\d+)(\s+h?)', expand = True)

weratedogs_archive_clean['name'] = weratedogs_archive_clean.text.str.extract('(is\s[A-Z][a-z]+\.)', expand = True)
```

Test

In [1]:

```
weratedogs_archive_clean
```

Quality:

Define: Cleaning the data re-generated for the `name` and `rating_denominator` columns.

Code

In [23]:

```
#cleaning the data re-gathered and testing to see the results
weratedogs_archive_clean.name = weratedogs_archive_clean.name.str[3:-1]
weratedogs_archive_clean.rating_denominator = weratedogs_archive_clean.rating_denominator.str[1:]
```

Test

In [2]:

```
weratedogs_archive_clean
```

Quality:

Define: Merging all dataframes into one using the `.merge()` method on the `tweet_id` column.

Code

In [25]:

```
weratedogs_archive_clean = pd.merge(weratedogs_archive_clean,tweets_clean, on = 'tweet_id', how = 'left')
```

In [26]:

```
weratedogs_archive_clean = pd.merge(weratedogs_archive_clean,image_prediction_clean, on = 'tweet_id', how = 'left')
```

Test

In [27]:

```
#A check to see if the merge was successful  
weratedogs_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 2356 entries, 0 to 2355  
Data columns (total 26 columns):  
tweet_id          2356 non-null int64  
timestamp         2356 non-null object  
source            2356 non-null object  
text              2356 non-null object  
rating_numerator  1927 non-null object  
rating_denominator 2281 non-null object  
name              1127 non-null object  
doggo             98 non-null object  
pupper            272 non-null object  
puppo             37 non-null object  
blep              1 non-null object  
floof             23 non-null object  
snoot             0 non-null object  
retweets          2354 non-null float64  
number_of_likes   2354 non-null float64  
jpg_url           2075 non-null object  
img_num           2075 non-null float64  
p1                2075 non-null object  
p1_conf           2075 non-null float64  
p1_dog            2075 non-null object  
p2                2075 non-null object  
p2_conf           2075 non-null float64  
p2_dog            2075 non-null object  
p3                2075 non-null object  
p3_conf           2075 non-null float64  
p3_dog            2075 non-null object  
dtypes: float64(6), int64(1), object(19)  
memory usage: 497.0+ KB
```

Tidiness:

Define: doggo , floof , pupper , puppo , blep and snoot - I will be collapsing these columns into a stage column using the `.melt()` method.

Code

In [28]:

```
#collapsing the stage variable columns to effect proper structure
weratedogs_archive_clean = pd.melt(weratedogs_archive_clean, id_vars=['tweet_id', 'times
tamp', 'source', 'rating_numerator', 'rating_denominator',
                                'name', 'retweets'
                                , 'number_of_likes', 'text',
                                'jpg_url', 'img_nu
m', 'p1', 'p1_conf', 'p1_dog', 'p2', 'p2_conf', 'p2_dog', 'p3', 'p3_conf', 'p3_dog'],
                                var_name='header', value_name = 'stage')

weratedogs_archive_clean = weratedogs_archive_clean.drop('header', axis =1)
```

Test

In [3]:

```
#checking if code implementation was a success
weratedogs_archive_clean
```

Quality:

Define: I will be tackling the duplicated rows created by the `.melt()` method used in the previous cell using the

`.drop_duplicates()` and `.drop()` methods.

Code

In [30]:

```
weratedogs_archive_clean.drop_duplicates(inplace = True)
```

In [31]:

```
#querying the duplicated tweet_id with Null values and leaving those with valid entries
mask_null_stage = weratedogs_archive_clean[weratedogs_archive_clean.tweet_id.duplicated
(keep = False)]
rows_to_drop = list(mask_null_stage[mask_null_stage.stage.isna()].index)
weratedogs_archive_clean = weratedogs_archive_clean.drop(rows_to_drop)
```

Test

In [32]:

```
#this test shows that we still have some form of duplicates hiding in our dataset  
weratedogs_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 2368 entries, 0 to 11021  
Data columns (total 21 columns):  
tweet_id          2368 non-null int64  
timestamp         2368 non-null object  
source            2368 non-null object  
rating_numerator  1938 non-null object  
rating_denominator 2293 non-null object  
name              1129 non-null object  
retweets          2366 non-null float64  
number_of_likes   2366 non-null float64  
text              2368 non-null object  
jpg_url           2087 non-null object  
img_num           2087 non-null float64  
p1                2087 non-null object  
p1_conf           2087 non-null float64  
p1_dog            2087 non-null object  
p2                2087 non-null object  
p2_conf           2087 non-null float64  
p2_dog            2087 non-null object  
p3                2087 non-null object  
p3_conf           2087 non-null float64  
p3_dog            2087 non-null object  
stage             431 non-null object  
dtypes: float64(6), int64(1), object(14)  
memory usage: 407.0+ KB
```

Quality:

Define: I will be tackling the duplicated `tweet_id` values which exposed double entry in the `stage` column using the

`.duplicated().` and `.drop()` methods.

Code

In [7]:

```
#checking if we still have duplicated tweet_ids  
weratedogs_archive_clean[weratedogs_archive_clean.tweet_id.duplicated(keep= False)]
```

In [34]:

```
#querying and deleting tweet_ids that have multiple stage entries as found in the cell above  
double_stage_entry = list(weratedogs_archive_clean[weratedogs_archive_clean.tweet_id.duplicated(keep= False)].index)  
weratedogs_archive_clean = weratedogs_archive_clean.drop(double_stage_entry)
```

Test


In [35]:

```
weratedogs_archive_clean[weratedogs_archive_clean.tweet_id.duplicated(keep= False)]
```

Out[35]:

tweet_id	timestamp	source	rating_numerator	rating_denominator	name	retweets	number_of_likes
----------	-----------	--------	------------------	--------------------	------	----------	-----------------

0 rows × 21 columns



Quality:

Define: Fiding and removing rows with missing image URLs from weratedogs_archive_clean dataframe.

Code

In [36]:

```
missing_image_Urls = list(weratedogs_archive_clean[weratedogs_archive_clean.jpg_url.isna()
a()].index)

weratedogs_archive_clean = weratedogs_archive_clean.drop(missing_image_Urls)
```

Test


In [37]:

```
weratedogs_archive_clean[weratedogs_archive_clean.jpg_url.isna()]
```

Out[37]:

tweet_id	timestamp	source	rating_numerator	rating_denominator	name	retweets	number_of_likes
----------	-----------	--------	------------------	--------------------	------	----------	-----------------

0 rows × 21 columns



Quality:

Define : Filling in the null values present in the retweets and number_of_likes columns using the .fillna method.

Code


In [38]:

```
#checking the rows with null values for retweets abd number_of_likes columns
weratedogs_archive_clean[weratedogs_archive_clean.number_of_likes.isna()]
```

Out[38]:

	tweet_id	timestamp	source	rating_numerator
19	888202515573088257	2017-07-21 01:02:36 +0000	<a href="http://twitter.com/download/iphone" r...	13
3171	771004394259247104	2016-08-31 15:19:06 +0000	<a href="http://twitter.com/download/iphone" r...	12

2 rows × 21 columns



In [39]:

```
#filling the missing values with the average number of retweets and Likes in the dataset
weratedogs_archive_clean['retweets'] = weratedogs_archive_clean['retweets'].fillna((weratedogs_archive_clean['retweets'].mean()))
weratedogs_archive_clean['number_of_likes'] = weratedogs_archive_clean['number_of_likes'].fillna((weratedogs_archive_clean['number_of_likes'].mean()))
```

Test

In [40]:

```
#checking if there are any null values left in the 'retweets' and 'number_of_likes' columns
weratedogs_archive_clean[weratedogs_archive_clean.number_of_likes.isna()].size, weratedogs_archive_clean[weratedogs_archive_clean.retweets.isna()].size
```

Out[40]:

(0, 0)

Quality:

Define : Converting all inappropriate data type to preferred data types.

Code

In [41]:

```
weratedogs_archive_clean.name = weratedogs_archive_clean.name.str.title()
weratedogs_archive_clean.p1_dog = weratedogs_archive_clean.p1_dog.astype(bool)
weratedogs_archive_clean.p2_dog = weratedogs_archive_clean.p2_dog.astype(bool)
weratedogs_archive_clean.p3_dog = weratedogs_archive_clean.p3_dog.astype(bool)
weratedogs_archive_clean.retweets= weratedogs_archive_clean.retweets.astype(int)
weratedogs_archive_clean.number_of_likes= weratedogs_archive_clean.number_of_likes.as
ype(int)
weratedogs_archive_clean.img_num=weratedogs_archive_clean.img_num.astype(int)
weratedogs_archive_clean.tweet_id = weratedogs_archive_clean.tweet_id.astype(str)
weratedogs_archive_clean.p1 = weratedogs_archive_clean.p1.str.title()
weratedogs_archive_clean.p2 = weratedogs_archive_clean.p2.str.title()
weratedogs_archive_clean.p3 = weratedogs_archive_clean.p3.str.title()
weratedogs_archive_clean.rating_numerator = weratedogs_archive_clean.rating_numerator.a
stype(float)
weratedogs_archive_clean.rating_denominator = weratedogs_archive_clean.rating_denominat
or.astype(float)
```

Test

In [42]:

```
weratedogs_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2063 entries, 0 to 11021
Data columns (total 21 columns):
tweet_id          2063 non-null object
timestamp         2063 non-null object
source            2063 non-null object
rating_numerator  1656 non-null float64
rating_denominator 2024 non-null float64
name              1044 non-null object
retweets          2063 non-null int64
number_of_likes   2063 non-null int64
text              2063 non-null object
jpg_url           2063 non-null object
img_num           2063 non-null int64
p1                2063 non-null object
p1_conf           2063 non-null float64
p1_dog            2063 non-null bool
p2                2063 non-null object
p2_conf           2063 non-null float64
p2_dog            2063 non-null bool
p3                2063 non-null object
p3_conf           2063 non-null float64
p3_dog            2063 non-null bool
stage             346 non-null object
dtypes: bool(3), float64(5), int64(3), object(10)
memory usage: 312.3+ KB
```

Quality:

Define : Filling in the null values present in the `rating_numerator` and `rating_denominator` columns using the `.fillna` method.

Code

In [8]:

```
#checking the rows with null values for rating_numerator column  
weratedogs_archive_clean[weratedogs_archive_clean.rating_numerator.isna()]
```

In [9]:

```
#checking the rows with null values for rating_denominator column  
weratedogs_archive_clean[weratedogs_archive_clean.rating_denominator.isna()]
```

In [45]:

```
#filling the missing values with the average number of rating_numerator and rating_denominator respectively  
  
weratedogs_archive_clean['rating_numerator'] = weratedogs_archive_clean['rating_numerator'].fillna(weratedogs_archive_clean['rating_numerator'].mean())  
weratedogs_archive_clean['rating_denominator'] = weratedogs_archive_clean['rating_denominator'].fillna(weratedogs_archive_clean['rating_denominator'].mean())
```

Test

In [46]:

```
weratedogs_archive_clean[weratedogs_archive_clean.rating_numerator.isna()].size,weratedogs_archive_clean[weratedogs_archive_clean.rating_denominator.isna()].size
```

Out[46]:

```
(0, 0)
```

Tidiness:

Define: Removing all rows with **'False'** dog predictions by the neural network.

Code

In [47]:

```
false_prediction_rows = list(weratedogs_archive_clean.query("p1_dog == False").index)
```

In [48]:

```
weratedogs_archive_clean = weratedogs_archive_clean.drop(false_prediction_rows, axis=0)
```

Test

In [49]:

```
weratedogs_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1521 entries, 1 to 11021
Data columns (total 21 columns):
tweet_id          1521 non-null object
timestamp         1521 non-null object
source            1521 non-null object
rating_numerator  1521 non-null float64
rating_denominator 1521 non-null float64
name              791 non-null object
retweets          1521 non-null int64
number_of_likes   1521 non-null int64
text              1521 non-null object
jpg_url           1521 non-null object
img_num           1521 non-null int64
p1                1521 non-null object
p1_conf           1521 non-null float64
p1_dog            1521 non-null bool
p2                1521 non-null object
p2_conf           1521 non-null float64
p2_dog            1521 non-null bool
p3                1521 non-null object
p3_conf           1521 non-null float64
p3_dog            1521 non-null bool
stage             255 non-null object
dtypes: bool(3), float64(5), int64(3), object(10)
memory usage: 230.2+ KB
```

Storing Data

Save gathered, assessed, and cleaned master dataset to a CSV file named "twitter_archive_master.csv".

In [50]:

```
weratedogs_archive_clean.to_csv(r'twitter_archive_master.csv', index=False)
```

Analyzing and Visualizing Data

In this section, analyze and visualize your wrangled data. You must produce at least **three (3) insights and one (1) visualization**.

In [51]:

```
df = pd.read_csv('twitter_archive_master.csv')
```

In [10]:

```
#Top viral tweets from the dataframe  
df.sort_values(by=['retweets'], ascending = False).head(12)
```

In [11]:

```
# Top 10 Tweets with the highest Likes  
df.sort_values(by=['number_of_likes'], ascending = False).head(10)
```

In [54]:

```
#The top five most frequent dog breed predicted by the neural network  
df.p1.value_counts().head()
```

Out[54]:

```
Golden_Retriever      143  
Labrador_Retriever    99  
Pembroke              89  
Chihuahua             83  
Pug                   57  
Name: p1, dtype: int64
```

Insights:


1. The top ten dog rating tweets with the highest retweets(coverage) in the dataset. On the top of the chart is a dog in its 'doggo' stage
2. The top ten most admired dog rating tweets. On the top of the chart is a dog in its 'puppo' stage
3. The top five most popular dog breeds in the neural network prediction.

Visualization

In [55]:

```
stages = df[~(df.stage.isna())]

plt.figure(figsize=(7,7))
plt.hist(stages.stage)
plt.title("Most Popular Dog Stage")
plt.ylabel('Frequency')
plt.xlabel('Stages');
```



In []: